

From Surface Realism to Behavioral Validity: Rethinking Evaluation for Repurposed Interactive LLM Systems

Anonymous ACL submission

Abstract

Large language models (LLMs) are increasingly repurposed as interactive systems, including counselor agents, simulated clients, and patient agents. Yet their evaluation often remains closer to static, output-centered benchmarking than to the behavioral and human-facing claims attached to these systems. We examine this mismatch through a structured review of 52 LLM-based counseling, therapy, and behavioral simulation papers. Our analysis shows that evaluation is abundant but often local: papers frequently use automatic metrics, expert ratings, LLM judges, rubrics, and domain-aware criteria, yet less often test behavioral trajectories, intervention-sensitive dynamics, evaluator validity, or consequences for users. We also find that theory-related constructs commonly reach evaluation, but measurement remains uneven, often relying on paper-specific rubrics, simulated outcomes, or local quality judgments rather than standardized or claim-matched evidence. Building on this diagnosis, we propose behavior-grounded evaluation for repurposed interactive LLM systems. This framework links domain constructs to observable behavioral markers, tracks behavior over interaction, validates evaluator ecosystems, and matches evidence to the claims being made. We argue that NLP progress for repurposed interactive LLM systems should not be defined only by higher local scores, but by whether systems exhibit the intended behavior under meaningful interactional conditions with evidence appropriate to their human-facing use.

1 Introduction

Large language models (LLMs) are changing what it means to evaluate progress in NLP. For much of the field’s history, progress has been organized around static benchmarks: fixed datasets, shared tasks, reference labels, and leaderboard metrics. This paradigm has been powerful because it makes model comparison tractable and cumulative. How-

ever, many current LLM systems are no longer used only as predictors over static inputs. General-purpose models are increasingly repurposed as interactive agents: counselors, simulated clients, patient agents, trainee partners, emotional-support systems, and role-play companions (Wang et al., 2024b, 2025b; Lee et al., 2024; Sungarda et al., 2025). In these settings, progress cannot be judged only by whether an output is fluent, preferred, or locally appropriate. It also depends on whether the system exhibits the behavior implied by its role, maintains that behavior across interaction, responds meaningfully to user actions, and supports the human-facing claims made about the system.

This creates a field-level evaluation mismatch. A simulated client may produce plausible utterances while drifting from its profile or responding inconsistently to counseling strategies (Yang et al., 2025b; Wang et al., 2026); a counseling agent may score highly on response quality while failing to maintain therapeutic stance or follow up on risk (Gabriel et al., 2024; Cai et al., 2026). The issue is therefore not that current work lacks evaluation. Rather, evaluation often remains local: it measures outputs, rubrics, and simulated interactions more readily than behavioral trajectories, intervention-sensitive dynamics, evaluator validity, or human-facing consequences.

LLM-based counseling, therapy, and role-play simulation provide a useful stress case for this mismatch. These systems are often motivated by goals such as counseling support, trainee practice, therapeutic usefulness, safety, or client simulation. They also rely on clinical, counseling, psychological, and behavioral constructs such as empathy, alliance, therapeutic stance, resistance, and risk handling (Li et al., 2024; Kim et al., 2025c). Many of these constructs are interactional: they unfold across turns, depend on user actions, and are interpreted relative to domain theory and use context. Recent reviews of mental-health LLMs and conversational agents

highlight concerns around reliability, clinical applicability, safety, user experience, standardized evaluation, and longitudinal evidence (Hua et al., 2025; Bucher et al., 2025; Shao et al., 2025). However, existing reviews do not specifically diagnose how current NLP and HCI papers evaluate repurposed interactive LLM systems, or what kinds of evidence are used to support claims about behavior, training value, safety, and human-facing use.

We address this gap through a structured review of 52 LLM-based counseling, therapy, role-play, and behavioral simulation papers. Our goal is not to provide an exhaustive review of digital mental-health systems. Instead, we use this domain as an empirical case for a broader question: how should NLP evaluate progress when general-purpose models are reused as interactive systems? We analyze each paper along dimensions covering system design, evaluation authority, theory grounding, trajectory evidence, and human-facing outcomes.

Our analysis shows that evaluation in this area is abundant but unevenly aligned with system claims. Papers frequently use automatic metrics, expert ratings, LLM judges, rubrics, and domain-aware criteria (Xiao et al., 2024; Feng et al., 2025; Xu et al., 2025b). Yet these evaluations often provide evidence about local output acceptability rather than behavioral validity. We identify five recurring gaps: evaluation authority is broad but unevenly grounded; output quality is overused as a proxy for behavioral validity; theory reaches evaluation but measurement remains uneven; dynamic design outpaces trajectory evaluation; and human-facing motivations often outpace human-facing evidence.

Building on this diagnosis, we propose *behavior-grounded evaluation* for repurposed interactive LLM systems. Behavior-grounded evaluation asks whether a system exhibits the intended behavior over time, under meaningful interactional conditions, with evidence appropriate to the claim. It links domain constructs to observable behavioral markers, tests those markers across trajectories, probes intervention sensitivity, matches evidence to human-facing claims, validates evaluator ecosystems, and packages interaction protocols for reproducibility.

This paper makes three contributions. First, it provides a structured empirical diagnosis of evaluation practice in 52 LLM-based counseling, therapy, role-play, and behavioral simulation papers. Second, it identifies recurring gaps between current evaluation practices and the claims made about

repurposed interactive LLM systems. Third, it articulates behavior-grounded evaluation as a claim-matched agenda for assessing interactive systems whose value depends on behavior over time, not only output quality. Together, these contributions argue that NLP progress in the LLM era should not be defined only by higher benchmark scores, but by whether evaluation can support the behavioral and human-facing claims attached to increasingly interactive systems.

2 Background: From Static Benchmarks to Repurposed Interactive LLM Systems

Prior reviews have established the rapid growth and difficulty of applying LLMs and conversational agents in mental-health contexts. Broad scoping reviews of mental-health LLMs highlight recurring concerns around reliability, clinical applicability, safety, and evaluation design (Hua et al., 2025). Integrative reviews of mental-health conversational agents further show that computer-science work often emphasizes model capability and response quality, whereas health-oriented work more often focuses on user outcomes and clinical relevance (Cho et al., 2023). Work on LLM-based medical-patient communication similarly emphasizes the gap between benchmark-centered model development and the interactional demands of clinical communication, where scenario alignment, communication quality, and human-centered evaluation become central (Shao et al., 2025). Together, these reviews motivate the importance of the domain, but they mainly characterize application areas, risks, capabilities, or broad evaluation challenges.

Our focus is different. We examine evaluation practice itself: who judges model behavior, what is measured, how theory is operationalized, whether dynamic behavior is evaluated over time, and whether evidence matches the human-facing claims made about these systems. This focus matters because LLM-based counseling and role-play systems are not only text generators. They are repurposed interactive systems: general-purpose models adapted through prompts, roles, memory, scaffolds, or interaction protocols to act as counselors, simulated clients, patient agents, or trainee partners. Their value depends not only on local response quality, but on whether they maintain coherent behavior, respond appropriately to user actions, and support claims about training, therapeutic use-

fulness, or simulation validity.

This concern is also visible in reviews focused on training, psychotherapy, and user experience. Reviews of conversational agents for mental-health training note challenges around standardized evaluation, authentic emotional responses, ethical safeguards, and clinical efficacy (Atapattu et al., 2025). Reviews of psychotherapy-oriented LLMs emphasize that assessment and intervention are dynamic and contextual, requiring systems that can adapt across stages of interaction and remain grounded in psychological theory (Na et al., 2025). Work on LLM mental-health applications also points to underexplored issues around design context, safety, user experience, and longitudinal outcomes (Bucher et al., 2025). These concerns make counseling and role-play simulation a useful stress case for a broader NLP evaluation problem.

We do not aim to provide an exhaustive review of digital mental-health systems. Instead, we use counseling and role-play simulation as a stress case for a broader NLP evaluation problem: how current NLP and HCI papers evaluate repurposed interactive LLM systems, and what kinds of evidence are needed when models are reused for behavioral, educational, therapeutic, or safety-relevant purposes. Existing reviews establish that evaluation is a recurring challenge; our review analyzes how evaluation is currently being done.

3 Method: Structured Review and Coding

We constructed and coded a corpus of LLM-based counseling and therapy papers, including closely related role-play and behavioral simulation work when it addressed client, patient, counselor, or psychologically grounded interaction. The review focused on papers that evaluated interactive generation, simulation behavior, counseling behavior, user experience, safety, or model performance.

3.1 Corpus construction

We searched ACL Anthology, ACM Digital Library, Google Scholar, and targeted HCI journals including IJHCS and IJHCI to cover NLP, HCI, and adjacent digital-health venues. Search queries combined terms for LLMs and generative agents with terms for counseling, psychotherapy, mental health, patient/client simulation, and evaluation. The initial search returned 900 records. After title/abstract screening, duplicate removal, and snow-




Figure 1: Corpus construction and screening process.

balling from eight highly relevant seed papers, we retained 112 candidate papers. Full-text relevance screening produced a final corpus of 52 primary papers. Figure 1 summarizes the corpus construction and screening process.

3.2 Inclusion criteria

We included papers that met three criteria: (1) they studied LLM-based or generative-agent systems; (2) they involved counseling, therapy, mental-health support, patient/client simulation, or behavioral/personality simulation; and (3) they reported some form of evaluation relevant to interaction, simulation behavior, counseling behavior, user experience, safety, or model performance. We excluded papers focused only on mental-health classification or detection, papers without interactive generation or simulation, papers outside the counseling simulation scope, and papers without accessible full text.

3.3 Coding scheme and validation

We iteratively developed a codebook with 38 fields covering simulation design, evaluation methods, human-centered targets, theory grounding, reliability/transparency, and robustness. Key fields included dynamic state, interaction level, longitudinal evaluation, evaluator type, LLM-judge use and validation, rubric use and reliability, theory grounding, theory operationalization, safety, realism, utility, fidelity, consistency, and human learning or outcome evidence.

Two student researchers coded the papers with ambiguous cases resolved through team discussion with the faculty lead. For fields central to our claims, we conducted additional audit passes to ensure that labels were supported by textual evidence.

4 Field Diagnosis: What Evaluation Measures and Misses

This section reports the main evaluation patterns in the 52-paper corpus and diagnoses what they leave under-tested. We organize the findings around five questions: who judges model behavior, what is measured, how theory enters evaluation, whether interactive behavior is evaluated over time, and whether evidence reaches the human-facing purposes used to motivate these systems. Across these dimensions, the central pattern is not an absence of evaluation. Rather, evaluation is abundant but often local: it measures outputs, rubrics, and simulated interactions more often than trajectories, consequences, or claim-specific evidence.

4.1 Evaluation authority is broad but unevenly grounded

The corpus shows a shift from single-metric evaluation toward evaluator triangulation. Automatic metrics appeared in 40 of 52 papers, human experts in 39, and LLM judges in 30; 21 papers combined all three. Examples include studies that combine expert or clinician ratings with LLM-based rubric scoring for counseling quality, safety, or therapeutic appropriateness (Qiu et al., 2025; Aghakhani et al., 2025; Cai et al., 2026), as well as studies that pair task-specific automatic metrics with human or LLM judgments for CBT, MI, or emotional-support behaviors (Xiao et al., 2024; Xu et al., 2025a; Gabriel et al., 2024). This is more informative than relying on automatic metrics alone, but it also requires clearer boundaries: automatic met-

rics, expert ratings, and LLM judges each support different kinds of evidence, and their scores should not be treated as interchangeable.

LLM-as-judge evaluation was common: 30 papers used LLM judges, and 26 reported some form of validation. This suggests that the field is not simply replacing human evaluation with unchecked LLM scores. The remaining issue is narrower but important: agreement with human raters can validate an LLM judge as a proxy scorer, but it does not by itself validate the rubric or the broader claim made from the score.

Rubric-based evaluation also creates a related issue. Forty-four papers used rubrics, but only 23 reported reliability or agreement (e.g., Cohen's kappa) for rubric-based judgments (Yang et al., 2025b; Feng et al., 2025). This matters because rubrics often define the central constructs being evaluated, such as empathy, safety, therapeutic appropriateness, or theory consistency. When reliability is not reported, it is difficult to know whether scores reflect stable criteria or evaluator-specific interpretations.

Thus, evaluation authority in this corpus is broad but unevenly grounded. The field increasingly triangulates among experts, LLM judges, and automatic metrics, but triangulation alone is not validation. The evidential force of these judgments depends on whether evaluators, rubrics, and validation procedures are aligned with the construct being measured and the claim being made.

4.2 Output quality is overused as a proxy for behavioral validity

The corpus shows substantial evaluation activity, but much of it remains close to the output surface. All 52 papers evaluated at least one local quality dimension, and 39 evaluated three or more. Utility was the most common target (46 papers), followed by fidelity (35), realism (30), emotional plausibility (30), consistency (22), and safety (22). These dimensions are necessary: unsafe, implausible, or task-unfaithful responses are clearly inadequate. The problem is that they mainly establish local acceptability.

Across papers, local acceptability is operationalized through measures such as empathy, coherence, guidance, authenticity, emotional expressiveness, active listening, relationship building, or therapy-technique use (Deshpande et al., 2025; Xie et al., 2025; Basar et al., 2025; Xu et al., 2025b). These measures can show that a response or short ex-

change looks appropriate. They provide weaker evidence about the system’s behavior as an interactive agent, where validity depends on patterns that unfold across turns rather than on isolated judgments of local acceptability. In this sense, the field is not failing to evaluate; it is often asking output-level measures to stand in for behavioral evidence.

4.3 Theory reaches evaluation, but measurement remains uneven

Theory is not absent from this corpus; the sharper issue is how theory becomes measurement. Many papers incorporate clinical, counseling, psychological, or behavioral constructs into their framing, design, or evaluation. Forty of 52 papers showed strong theory grounding, and another 10 showed partial grounding. Theory-related constructs also appeared frequently in evaluation: 30 papers strongly operationalized such constructs, 16 did so partially, and 6 showed no clear theory-shaped evaluation.

The strongest cases made theory measurable through established psychological scales, therapy coding systems, theory-specific benchmarks, and construct-specific rubrics. Examples include therapeutic alliance and counseling-skill measures in CARE-Bench (Wang et al., 2026) and Client-CAST (Wang et al., 2024a), motivational interviewing criteria in KMI (Kim et al., 2025a) and CAMI (Yang et al., 2025a), and CBT-oriented ratings in CACTUS (Lee et al., 2024). Across the corpus, however, standardized instruments were less common than locally defined evaluation schemes: 9 papers used validated scales and 7 used established therapy or counseling coding systems, while 35 used theory-specific rubrics and 26 used theory-specific benchmark tasks. This suggests that the field often translates theory into evaluable dimensions, but not always through shared or standardized measurement instruments.

Even established instruments (e.g., PANAS, WAI) require caution in this setting. Measures of alliance, affect, empathy, symptoms, or counseling competence were originally designed for specific human respondents, trained observers, or clinical contexts. When such measures are answered by LLM-simulated clients or scored by LLM judges, the same score changes meaning: it may indicate simulated state consistency or agreement with local ratings, but not actual affective change, therapeutic alliance, user benefit, or safety. Thus, a recognized scale or rubric is not automatically strong evidence.

It matters who produces the score and what claim the score is used to support.

The partial cases further show the unevenness of theory-shaped measurement. Sixteen papers relied on broader or paper-specific measures, including realism, authenticity, profile adherence, client resemblance, counseling quality, training usefulness, diagnostic correctness, or interaction quality (e.g., PsyCLIENT (Qiu et al., 2026), Roleplaydoh (Louie et al., 2024)). These measures are relevant to mental-health and counseling applications, but they are often defined differently across papers. As a result, it is difficult to compare results across studies or to know whether similarly named constructs are being measured in the same way. Thus, the gap is less the absence of theory-shaped evaluation than the unevenness of its measurement. Theory-related constructs are often evaluated, but not always in ways that are standardized, comparable, or valid for the specific evaluator, respondent, or claim being supported.

4.4 Dynamic design outpaces trajectory evaluation

For interactive counseling and role-play systems, behavioral validity depends on how behavior unfolds over time. A plausible multi-turn dialogue is not necessarily a valid behavioral trajectory: the system must maintain coherent profiles, states, and response patterns as the interaction develops. This matters for systems that claim memory, adaptation, client-state modeling, therapeutic progression, or sensitivity to counselor interventions.

Yet dynamic design outpaces dynamic evaluation. Twenty-seven of 52 papers studied extended dialogues, and 19 used some form of dynamic state, such as memory, evolving profiles, state transitions, session summaries, cognitive diagrams, or emotion trajectories. Only 7 papers reported longitudinal evaluation; among the 19 papers with dynamic state, 13 did not evaluate behavior longitudinally. Similarly, 39 papers involved intervention-sensitive evaluation or design, but only 7 of these included longitudinal evaluation. Thus, many papers examine whether a system can respond plausibly within an interaction, but fewer test whether those responses form stable, coherent, and theoretically meaningful behavioral trajectories.

A few papers point toward more trajectory-aware evaluation. EmoAgent measures pre–post changes in simulated psychological assessments (Qiu et al., 2025); the Interactive Narrative Therapist work

evaluates therapeutic progression through innovative moments (Feng et al., 2025); and MusPsy uses multi-session counseling with working-alliance and affect measures across sessions (Wang et al., 2025a). These examples show that evaluation can track change across an interaction, even when the outcomes are simulated or proxy-based. However, they remain exceptions: dynamic behavior is often claimed or engineered, while evaluation still tends to reduce interaction to local judgments, aggregate scores, or single endpoint summaries.

4.5 Human-facing motivations often outpace human-facing evidence

The previous sections show that many papers evaluate outputs with domain-aware criteria. The remaining gap is whether such criteria are connected to the human-facing purposes used to motivate these systems, such as counseling support, trainee practice, therapeutic usefulness, or safety. User studies appeared in 19 of 52 papers, lay-user evaluation in 16, and explicit human learning or outcome measures in only 10. Across the corpus, 36 papers showed no clear evaluation beyond model outputs, benchmark performance, simulated dialogue quality, expert ratings, or LLM-judge scores.

This creates a claim–evidence gap. Utility was evaluated in 46 papers, but human learning or outcome evidence appeared in only 10. Safety was evaluated in 22 papers, but only a small number examined safety or deployment consequences. Similarly, 8 papers measured outcomes only in simulated clients or patients, such as simulated affect, alliance, symptom severity, or therapeutic progression. Such evidence is useful for testing synthetic interactions, but it should not be treated as evidence of real user benefit.

A few papers do connect evaluation to human-facing effects. Training-oriented studies measure counselor skill development (Louie et al., 2026; Tanana et al., 2019); perception-oriented studies measure perceived empathy or interaction experience (Schmidmaier et al., 2025b,a); and safety-oriented work evaluates high-risk mental-health dialogues (Cai et al., 2026). These examples remain exceptions. The core issue is not that every paper must evaluate all downstream effects, but that evidence should match the level of the claim: claims about training value, counseling support, therapeutic usefulness, or safety require more than local output quality or simulated outcomes.

5 Toward Behavior-Grounded Evaluation for Repurposed Interactive LLM Systems

The diagnosis above suggests that evaluation for repurposed interactive LLM systems should move from output-centered scoring to behavior-grounded evidence. We use *behavior-grounded evaluation* to refer to a layered evaluation logic: first define the domain construct being claimed, then identify observable behavioral markers, test whether those markers persist and change across interaction trajectories, probe whether they respond to meaningful interventions, and match the resulting evidence to the human-facing claim being made. This is not a call for every paper to run clinical trials or user studies. Rather, it requires authors to make explicit what level of claim their evidence supports. Output-quality claims, behavioral-simulation claims, training claims, safety claims, and deployment claims require different forms of evidence. The key shift is from asking whether an output looks good to asking whether a system exhibits the right behavior, under the right conditions, with evidence appropriate to the claim.

5.1 Construct-grounded evaluation

Evaluation should begin with the construct that a system claims to support or simulate. For repurposed interactive systems, broad labels such as “helpful”, “realistic”, or “therapeutic” are too underspecified to support strong claims. A counseling agent, for example, may need to demonstrate empathy, alliance-building, motivational interviewing adherence, CBT competence, or safety handling; a client simulator may need to represent avoidance, resistance, disclosure, emotional reactivity, or readiness to change. These constructs should be named, defined, and linked to the system claim before metrics are selected.

Construct-grounded evaluation therefore asks: what construct is the system expected to exhibit, and what would count as evidence for it? This does not require every study to use a validated clinical scale, but it does require making the evaluation target explicit. The key requirement is claim–construct alignment: the construct being measured should be the construct needed to support the paper’s claim. Without such grounding, high realism or usefulness scores may obscure whether the system actually exhibits the intended behavior.

5.2 Behavioral marker evaluation

Constructs should be translated into behavioral markers: observable dialogue behaviors that indicate whether an abstract construct is being expressed. Empathy, for example, can be marked by validation, reflection, or perspective-taking; safety can be marked by risk assessment, boundary setting, or escalation behavior. This translation is the bridge between construct labels and dialogue evidence: it turns theory into inspectable behavior rather than leaving it as a label.

Behavioral marker evaluation makes evaluation more diagnostic. Instead of asking whether a dialogue is generally “good”, “realistic”, or “therapeutic”, evaluators can ask whether expected behaviors occur, whether inappropriate behaviors are avoided, and whether the behavior matches the intended role. Without such markers, a dialogue can look plausible while still failing to exhibit the behavior the system is supposed to simulate or support.

5.3 Trajectory-based evaluation

Trajectory-based evaluation should define and track behavioral markers across an interaction. For systems that claim memory, adaptation, client-state modeling, or therapeutic progression, evaluation should specify how the relevant behavior is expected to persist, drift, or change as the interaction unfolds. A trajectory is therefore not simply a longer dialogue; it is an ordered behavioral path whose stability and development can be tested.

For a client simulator, evaluation should test whether the intended profile is preserved while relevant states change coherently across counselor actions, such as increased disclosure after empathic reflection or increased resistance after confrontation. For a counselor agent, evaluation should test whether the system maintains therapeutic stance, follows up on earlier disclosures, repairs rupture, handles risk consistently, or adapts intervention strategy over time. The unit of evaluation is not the isolated response, but the behavioral path that the system produces under interaction.

5.4 Intervention-sensitive evaluation

Intervention-sensitive evaluation should test behavioral contingencies, not merely response diversity. The evaluation should hold the client profile or scenario constant, vary the counselor action, and pre-specify the expected behavioral effect of each intervention. For example, with the same avoidant client

profile, empathic reflection might be expected to increase disclosure, confrontation might increase resistance, and cognitive reframing might shift interpretation of the problem. A behaviorally valid simulator should show these differentiated response patterns, rather than producing equally plausible continuations across all counselor actions.

This requires evaluating the effect of an intervention on subsequent markers, not only the immediate next response. For instance, after a risk disclosure, the system should show risk-assessment, boundary-setting, or escalation markers across the following turns, rather than generic reassurance. Thus, intervention-sensitive evaluation moves from asking whether a response sounds appropriate to asking whether meaningful actions produce theoretically expected behavioral consequences.

5.5 Claim-matched consequence evaluation

Claim-matched consequence evaluation links the strength of a claim to the level of evidence used to support it. Output-quality evidence can support claims about response generation or construct alignment; trajectory evidence can support claims about behavioral simulation; human-facing evidence is needed for claims about learning, support, or deployment readiness.

This distinction matters because evidence can be easily overextended. Simulated client outcomes can reveal synthetic interaction dynamics, but not real user benefit. Expert ratings can support clinical plausibility, but not necessarily trainee learning. LLM-judge scores can support scalable rubric scoring, but not user experience or safety in use. Behavior-grounded evaluation therefore requires making the boundary of evidence explicit: what the evaluation supports, what it does not support, and what stronger claims would require.

5.6 Validated evaluator ecosystems

Because repurposed interactive systems make different kinds of claims, no single evaluator is sufficient for all purposes. Automatic metrics can support scalable comparison, expert ratings can assess domain-specific plausibility, LLM judges can apply rubrics at scale, and users can report experience or perceived value. These sources of evidence should be treated as complementary rather than interchangeable.

A validated evaluator ecosystem specifies what each evaluator is trusted to judge. LLM judges should be calibrated against human or expert rat-

ings for the specific rubric they apply, and human rubrics should report reliability when multiple raters are involved. More broadly, validation should operate at three levels: whether the evaluator can apply the rubric reliably, whether the rubric captures the intended construct, and whether the resulting score supports the claim being made. Triangulation is valuable, but triangulation alone is not validation; the evaluator must match the construct and claim.

6 Implications for NLP Progress

Our case study suggests that the question of progress in NLP changes when general-purpose models are repurposed as interactive systems. In static benchmark settings, progress can often be framed as better performance on a shared task. In reused human-facing systems, however, the meaning of progress depends on the claim being made: better response generation, valid behavioral simulation, training value, safety, therapeutic usefulness, or deployment readiness require different forms of evidence.

First, progress should be claim-relative. A higher local quality score is meaningful for claims about response generation, but it is weaker evidence for claims about behavioral simulation, training value, or human benefit. Our review shows that many papers evaluate outputs with domain-aware criteria, but evidence often remains local to model outputs, simulated interactions, or evaluator judgments. For repurposed interactive systems, progress should therefore be stated as a claim–evidence relation: what kind of improvement is claimed, and what kind of evidence would make that improvement meaningful?

Second, generalization should include behavioral transfer. For interactive LLM systems, generalization is not only performance on unseen examples or scenarios. It also concerns whether the system preserves the intended behavioral pattern across turns, users, interventions, and contexts. A client simulator that appears realistic in one dialogue but loses its profile under different counselor actions has not generalized behaviorally. A counseling agent that produces empathic responses but fails to maintain stance during rupture, resistance, or risk has not generalized interactionally.

Third, evaluation should become an ecosystem rather than a single score. Automatic metrics, expert ratings, LLM judges, rubrics, and user studies

each provide different forms of evidence. The challenge is not simply to add more evaluators, but to define what each evaluator is trusted to judge and how those judgments are validated. Evaluator agreement may support local scoring reliability, but it does not automatically establish construct validity, trajectory validity, or human-facing benefit.

Overall, our findings suggest that benchmark performance is only one form of progress. For repurposed interactive LLM systems, progress also means producing behavior that is construct-grounded, stable across interaction, sensitive to meaningful interventions, and supported by evidence appropriate to the intended human use. This does not replace traditional evaluation; it extends it to settings where models are no longer only solving tasks, but participating in consequential interactions.

7 Conclusion

As LLMs are increasingly repurposed as interactive systems, NLP evaluation must account for more than local output quality. Using LLM-based counseling, therapy, and role-play simulation as a stress case, we reviewed 52 papers and found that evaluation is abundant but often local: studies frequently use automatic metrics, expert ratings, LLM judges, rubrics, and domain-aware criteria, yet less often test behavioral trajectories, intervention-sensitive dynamics, evaluator validity, or human-facing consequences.

The central issue is therefore not simply a lack of evaluation, but a mismatch between evidence and claim. Output-level or simulated evidence can show that responses look plausible, safe, empathic, or realistic, but it cannot by itself establish profile stability, therapeutic stance, trainee learning, user benefit, or deployment readiness. We argued for behavior-grounded evaluation as a claim-matched logic that links constructs to behavioral markers, tests them across trajectories and interventions, validates evaluator ecosystems, and makes evidence boundaries explicit.

For repurposed interactive LLM systems, progress should not be defined only by higher local scores. It should depend on whether systems exhibit the intended behavior, under meaningful interactional conditions, with evidence appropriate to the claims being made.

8 Limitations

References

- Elham Aghakhani, Lu Wang, Karla T. Washington, George Demiris, Jina Huh-Yoo, and Rezvaneh Reza-pour. 2025. [From conversation to automation: Leveraging LLMs for problem-solving therapy analysis](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 25189–25207, Vienna, Austria. Association for Computational Linguistics.
- Thushari Atapattu, Menasha Thilakaratne, Duc Nhan Do, Mahen Herath, and Katrina E. Falkner. 2025. [Exploring the role of mental health conversational agents in training medical students and professionals: A systematic literature review](#). In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 20785–20798, Vienna, Austria. Association for Computational Linguistics.
- Erkan Basar, Iris Hendrickx, Emiel Krahmer, Gert-Jan Bruijn, and Tibor Bosse. 2024. [To what extent are large language models capable of generating substantial reflections for motivational interviewing counseling chatbots? a human evaluation](#). In *Proceedings of the 1st Human-Centered Large Language Modeling Workshop*, pages 41–52, TBD. ACL.
- Erkan Basar, Xin Sun, Iris Hendrickx, Jan de Wit, Tibor Bosse, Gert-Jan De Bruijn, Jos A. Bosch, and Emiel Krahmer. 2025. [How well can large language models reflect? a human evaluation of LLM-generated reflections for motivational interviewing dialogues](#). In *Proceedings of the 31st International Conference on Computational Linguistics*, pages 1964–1982, Abu Dhabi, UAE. Association for Computational Linguistics.
- Santiago Berrezueta-Guzman, Mohanad Kandil, María-Luisa Martín-Ruiz, Iván Pau de la Cruz, and Stephan Krusche. 2024. [Future of adhd care: Evaluating the efficacy of chatgpt in therapy enhancement](#). *Health-care*, 12(6).
- Andreas Bucher, Sarah Egger, Inna Vashkite, Wenyan Wu, and Gerhard Schwabe. 2025. [“it’s not only attention we need”: Systematic review of large language models in mental health care](#). *JMIR Ment Health*, 12:e78410.
- Yunna Cai, Fan Wang, Haowei Wang, Kun Wang, Kailai Yang, Sophia Ananiadou, Moyan Li, and Mingming Fan. 2026. [Exploring safety alignment evaluation of llms in chinese mental health dialogues via llm-as-judge](#). *Preprint*, arXiv:2508.08236.
- Siyuan Chen, Mengyue Wu, Kenny Q. Zhu, Kunyao Lan, Zhiling Zhang, and Lychun Cui. 2023. [Llm-empowered chatbots for psychiatrist and patient simulation: Application and evaluation](#). *Preprint*, arXiv:2305.13614.
- Yujia Chen, Changsong Li, Yiming Wang, Tianjie Ju, Qingqing Xiao, Nan Zhang, Zifan Kong, Peng Wang, and Binyu Yan. 2025. [MIND: Towards immersive psychological healing with multi-agent inner dialogue](#). In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 9380–9413, Suzhou, China. Association for Computational Linguistics.
- Yu Ying Chiu, Ashish Sharma, Inna Wanyin Lin, and Tim Althoff. 2024. [A computational framework for behavioral assessment of llm therapists](#). *Preprint*, arXiv:2401.00820.
- Young Min Cho, Sunny Rai, Lyle Ungar, João Sedoc, and Sharath Guntuku. 2023. [An integrative survey on mental health conversational agents to bridge computer science and medical perspectives](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 11346–11369, Singapore. Association for Computational Linguistics.
- Neha Pravin Deshpande, Stefan Hillmann, and Sebastian Möller. 2025. [Evaluating large language models for enhancing live chat therapy: A comparative study with psychotherapists](#). In *Proceedings of the 26th Annual Meeting of the Special Interest Group on Discourse and Dialogue*, pages 800–812, Avignon, France. Association for Computational Linguistics.
- Yi Feng, Jiaqi Wang, Wenxuan Zhang, Zhuang Chen, Shen Yutong, Xiyao Xiao, Minlie Huang, Liping Jing, and Jian Yu. 2025. [Reframe your life story: Interactive narrative therapist and innovative moment assessment with large language models](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 24484–24509, Suzhou, China. Association for Computational Linguistics.
- Aya E Fouda, Abdelrahman A Hassan, Radwa J Hanafy, and Mohammed E Fouda. 2026. [Psychiatrybench: a multi-task benchmark for llms in psychiatry](#). *npj Digital Medicine*.
- Saadiah Gabriel, Isha Puri, Xuhai Xu, Matteo Malgaroli, and Marzyeh Ghassemi. 2024. [Can AI relate: Testing large language model response for mental health support](#). In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 2206–2221, Miami, Florida, USA. Association for Computational Linguistics.
- Kilichbek Haydarov, Youssef Mohamed, Emilio Goldenhersch, Paul OCallaghan, Li-jia Li, and Mohamed Elhoseiny. 2025. [Towards AI-assisted psychotherapy: Emotion-guided generative interventions](#). In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 32736–32755, Suzhou, China. Association for Computational Linguistics.
- Jinpeng Hu, Tengpeng Dong, Gang Luo, Hui Ma, Peng Zou, Xiao Sun, Dan Guo, Xun Yang, and Meng Wang. 2025. [Psychollm: Enhancing llm for psychological understanding and evaluation](#). *IEEE Transactions on Computational Social Systems*, 12(2):539–551.

981	Chen. 2025. A survey of large language models in psychotherapy: Current landscape and future directions . In <i>Findings of the Association for Computational Linguistics: ACL 2025</i> , pages 7362–7376, Vienna, Austria. Association for Computational Linguistics.	1037
982		1038
983		1039
984		1040
985		1041
986		1042
987		1043
988	Viet Cuong Nguyen, Mohammad Taher, Dongwan Hong, Vinicius Konkolics Possobom, Vibha Thirunelayi Gopalakrishnan, Ekta Raj, Zihang Li, Heather J. Soled, Michael L. Birnbaum, Srijan Kumar, and Munmun De Choudhury. 2025. Do large language models align with core mental health counseling competencies? In <i>Findings of the Association for Computational Linguistics: NAACL 2025</i> , pages 7503–7526, Albuquerque, New Mexico. Association for Computational Linguistics.	1044
989		1045
990		1046
991		1047
992		1048
993		1049
994		1050
995		1051
996		1052
997		1053
998	Mithat Can Ozgun, Jiahuan Pei, Koen Hindriks, Lucia Donatelli, Qingzhi Liu, and Junxiao Wang. 2025. Trustworthy ai psychotherapy: Multi-agent llm workflow for counseling and explainable mental disorder diagnosis . In <i>Proceedings of the 34th ACM International Conference on Information and Knowledge Management, CIKM '25</i> , pages 2263–2272, New York, NY, USA. Association for Computing Machinery.	1054
999		1055
1000		1056
1001		1057
1002		1058
1003		1059
1004		1060
1005		1061
1006		1062
1007		1063
1008		1064
1009		1065
1010		1066
1011		1067
1012		1068
1013		1069
1014		1070
1015		1071
1016		1072
1017		1073
1018		1074
1019		1075
1020		1076
1021		1077
1022		1078
1023		1079
1024		1080
1025		1081
1026		1082
1027		1083
1028		1084
1029		1085
1030		1086
1031		1087
1032		1088
1033		1089
1034		1090
1035		1091
1036		1092
1037		1093
1038		1094
1039		1095
1040		1096
1041		1097
1042		1098
1043		1099
1044		1100
1045		1101
1046		1102
1047		1103
1048		1104
1049		1105
1050		1106
1051		1107
1052		1108
1053		1109
1054		1110
1055		1111
1056		1112
1057		1113
1058		1114
1059		1115
1060		1116
1061		1117
1062		1118
1063		1119
1064		1120
1065		1121
1066		1122
1067		1123
1068		1124
1069		1125
1070		1126
1071		1127
1072		1128
1073		1129
1074		1130
1075		1131
1076		1132
1077		1133
1078		1134
1079		1135
1080		1136
1081		1137
1082		1138
1083		1139
1084		1140
1085		1141
1086		1142
1087		1143
1088		1144
1089		1145
1090		1146
1091		1147
1092		1148
1093		1149
1094		1150
1095		1151
1096		1152
1097		1153
1098		1154
1099		1155
1100		1156
1101		1157
1102		1158
1103		1159
1104		1160
1105		1161
1106		1162
1107		1163
1108		1164
1109		1165
1110		1166
1111		1167
1112		1168
1113		1169
1114		1170
1115		1171
1116		1172
1117		1173
1118		1174
1119		1175
1120		1176
1121		1177
1122		1178
1123		1179
1124		1180
1125		1181
1126		1182
1127		1183
1128		1184
1129		1185
1130		1186
1131		1187
1132		1188
1133		1189
1134		1190
1135		1191
1136		1192
1137		1193
1138		1194
1139		1195
1140		1196
1141		1197
1142		1198
1143		1199
1144		1200
1145		1201
1146		1202
1147		1203
1148		1204
1149		1205
1150		1206
1151		1207
1152		1208
1153		1209
1154		1210
1155		1211
1156		1212
1157		1213
1158		1214
1159		1215
1160		1216
1161		1217
1162		1218
1163		1219
1164		1220
1165		1221
1166		1222
1167		1223
1168		1224
1169		1225
1170		1226
1171		1227
1172		1228
1173		1229
1174		1230
1175		1231
1176		1232
1177		1233
1178		1234
1179		1235
1180		1236
1181		1237
1182		1238
1183		1239
1184		1240
1185		1241
1186		1242
1187		1243
1188		1244
1189		1245
1190		1246
1191		1247
1192		1248
1193		1249
1194		1250
1195		1251
1196		1252
1197		1253
1198		1254
1199		1255
1200		1256
1201		1257
1202		1258
1203		1259
1204		1260
1205		1261
1206		1262
1207		1263
1208		1264
1209		1265
1210		1266
1211		1267
1212		1268
1213		1269
1214		1270
1215		1271
1216		1272
1217		1273
1218		1274
1219		1275
1220		1276
1221		1277
1222		1278
1223		1279
1224		1280
1225		1281
1226		1282
1227		1283
1228		1284
1229		1285
1230		1286
1231		1287
1232		1288
1233		1289
1234		1290
1235		1291
1236		1292
1237		1293
1238		1294
1239		1295
1240		1296
1241		1297
1242		1298
1243		1299
1244		1300
1245		1301
1246		1302
1247		1303
1248		1304
1249		1305
1250		1306
1251		1307
1252		1308
1253		1309
1254		1310
1255		1311
1256		1312
1257		1313
1258		1314
1259		1315
1260		1316
1261		1317
1262		1318
1263		1319
1264		1320
1265		1321
1266		1322
1267		1323
1268		1324
1269		1325
1270		1326
1271		1327
1272		1328
1273		1329
1274		1330
1275		1331
1276		1332
1277		1333
1278		1334
1279		1335
1280		1336
1281		1337
1282		1338
1283		1339
1284		1340
1285		1341
1286		1342
1287		1343
1288		1344
1289		1345
1290		1346
1291		1347
1292		1348
1293		1349
1294		1350
1295		1351
1296		1352
1297		1353
1298		1354
1299		1355
1300		1356
1301		1357
1302		1358
1303		1359
1304		1360
1305		1361
1306		1362
1307		1363
1308		1364
1309		1365
1310		1366
1311		1367
1312		1368
1313		1369
1314		1370
1315		1371
1316		1372
1317		1373
1318		1374
1319		1375
1320		1376
1321		1377
1322		1378
1323		1379
1324		1380
1325		1381
1326		1382
1327		1383
1328		1384
1329		1385
1330		1386
1331		1387
1332		1388
1333		1389
1334		1390
1335		1391
1336		1392
1337		1393
1338		1394
1339		1395
1340		1396
1341		1397
1342		139

1094	and Yifei Zhang. 2025b. AnnaAgent: Dynamic evolution agent system with multi-session memory for realistic seeker simulation . In <i>Findings of the Association for Computational Linguistics: ACL 2025</i> , pages 23221–23235, Vienna, Austria. Association for Computational Linguistics.	1152
1095		1153
1096		1154
1097		1155
1098		1156
1099		1157
1100	Ruiyi Wang, Stephanie Milani, Jamie C. Chiu, Jiayin Zhi, Shaun M. Eack, Travis Labrum, Samuel M Murphy, Nev Jones, Kate V Hardy, Hong Shen, Fei Fang, and Zhiyu Chen. 2024b. PATIENT-ψ: Using large language models to simulate patients for training mental health professionals . In <i>Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing</i> , pages 12772–12797, Miami, Florida, USA. Association for Computational Linguistics.	1158
1101		
1102		1159
1103		1160
1104		1161
1105		1162
1106		1163
1107		1164
1108		1165
1109		1166
1110	Xiaoyi Wang, Jiwei Zhang, Guangtao Zhang, and Honglei Guo. 2025c. Feel the difference? a comparative analysis of emotional arcs in real and LLM-generated CBT sessions . In <i>Findings of the Association for Computational Linguistics: EMNLP 2025</i> , pages 19999–20017, Suzhou, China. Association for Computational Linguistics.	1167
1111		
1112		1169
1113		1170
1114		1171
1115		1172
1116		1173
1117	Mengxi Xiao, Qianqian Xie, Ziyang Kuang, Zhicheng Liu, Kailai Yang, Min Peng, Weiguang Han, and Jimin Huang. 2024. HealMe: Harnessing cognitive reframing in large language models for psychotherapy . In <i>Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 1707–1725, Bangkok, Thailand. Association for Computational Linguistics.	1174
1118		1175
1119		1176
1120		1177
1121		
1122		
1123		
1124		
1125	Haojie Xie, Yirong Chen, Xiaofen Xing, Jingkai Lin, and Xiangmin Xu. 2025. PsyDT: Using LLMs to construct the digital twin of psychological counselor with personalized counseling style for psychological counseling . In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 1081–1115, Vienna, Austria. Association for Computational Linguistics.	
1126		
1127		
1128		
1129		
1130		
1131		
1132		
1133	Jia Xu, Tianyi Wei, Bojian Hou, Patryk Orzechowski, Shu Yang, Ruochen Jin, Rachael Paulbeck, Joost Wagenaar, George Demiris, and Li Shen. 2025a. Mentalchat16k: A benchmark dataset for conversational mental health assistance . In <i>Proceedings of the 31st ACM SIGKDD Conference on Knowledge Discovery and Data Mining V.2, KDD '25</i> , pages 5367–5378, New York, NY, USA. Association for Computing Machinery.	
1134		
1135		
1136		
1137		
1138		
1139		
1140		
1141		
1142	Yangyang Xu, Jinpeng Hu, Zhuoer Zhao, Zhangling Duan, Xiao Sun, and Xun Yang. 2025b. MultiAgentESC: A LLM-based multi-agent collaboration framework for emotional support conversation . In <i>Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing</i> , pages 4665–4681, Suzhou, China. Association for Computational Linguistics.	
1143		
1144		
1145		
1146		
1147		
1148		
1149		
1150	Yizhe Yang, Palakorn Achananuparp, Heyan Huang, Jing Jiang, Phey Ling Kit, Nicholas Gabriel Lim,	
1151		
	Cameron Tan Shi Ern, and Ee-Peng Lim. 2025a. CAMI: A counselor agent supporting motivational interviewing through state inference and topic exploration . In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 21037–21081, Vienna, Austria. Association for Computational Linguistics.	
	Yizhe Yang, Palakorn Achananuparp, Heyan Huang, Jing Jiang, Nicholas Gabriel Lim, Cameron Tan Shi Ern, Phey Ling Kit, Jenny Giam Xiuhui, John Pinto, and Ee-Peng Lim. 2025b. Consistent client simulation for motivational interviewing-based counseling . In <i>Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)</i> , pages 20959–20998, Vienna, Austria. Association for Computational Linguistics.	
	Mian Zhang, Xianjun Yang, Xinlu Zhang, Travis Labrum, Jamie C. Chiu, Shaun M. Eack, Fei Fang, William Yang Wang, and Zhiyu Chen. 2025. CBT-bench: Evaluating large language models on assisting cognitive behavior therapy . In <i>Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)</i> , pages 3864–3900, Albuquerque, New Mexico. Association for Computational Linguistics.	

Category	Strong	Partial	Mentioned	Not Specified	Total
Theory Grounding	40 (76.9%)	10 (19.2%)	2 (3.8%)	—	52
Theory Operationalization	30 (57.7%)	16 (30.8%)	—	6 (11.5%)	52

Table 1: Statistics of Theory Grounding and Theory Operationalization in Evaluation

Category	Papers
Strong	(Xiao et al., 2024), (Gabriel et al., 2024), (Haydarov et al., 2025), (Qiu et al., 2025), (Aghakhani et al., 2025), (Wang et al., 2026), (Sungarda et al., 2025), (Chiu et al., 2024), (Li et al., 2024), (Wang et al., 2024a), (Fouda et al., 2026), (Cai et al., 2026), (Nguyen et al., 2025), (Zhang et al., 2025), (Kim et al., 2025a), (Jaiswal et al., 2024), (Xu et al., 2025a), (Schmidmaier et al., 2025b), (Schmidmaier et al., 2025a), (Ozgun et al., 2025), (Louie et al., 2026), (Wang et al., 2024b), (Lee et al., 2024), (Basar et al., 2024), (Yang et al., 2025b), (Yang et al., 2025a), (Xie et al., 2025), (Basar et al., 2025), (Feng et al., 2025), (Xu et al., 2025b), (Kim et al., 2025c), (Kim et al., 2025b), (Liu et al., 2025), (Wang et al., 2025c), (Samuel et al., 2025), (Chen et al., 2025), (Na, 2024), (Sharma et al., 2023), (Wang et al., 2025a), (Tanana et al., 2019)
Partial	(Wang et al., 2025b), (Liu et al., 2023), (Hu et al., 2025), (Qiu et al., 2026), (Deshpande et al., 2025), (Louie et al., 2024), (Kermani et al., 2025), (Kang et al., 2024), (Chen et al., 2023), (Schmidgall et al., 2025)
Mentioned	(Kuhail et al., 2025), (Berrezueta-Guzman et al., 2024)

Table 2: Papers Classified by Theory Grounding Level

Category	Papers
Strong	(Xiao et al., 2024), (Gabriel et al., 2024), (Qiu et al., 2025), (Aghakhani et al., 2025), (Wang et al., 2026), (Sungarda et al., 2025), (Chiu et al., 2024), (Li et al., 2024), (Wang et al., 2024a), (Cai et al., 2026), (Nguyen et al., 2025), (Zhang et al., 2025), (Kim et al., 2025a), (Xu et al., 2025a), (Schmidmaier et al., 2025b), (Schmidmaier et al., 2025a), (Louie et al., 2026), (Wang et al., 2024b), (Lee et al., 2024), (Yang et al., 2025a), (Feng et al., 2025), (Kim et al., 2025c), (Kim et al., 2025b), (Wang et al., 2025c), (Samuel et al., 2025), (Na, 2024), (Sharma et al., 2023), (Kang et al., 2024), (Wang et al., 2025a), (Tanana et al., 2019)
Partial	(Haydarov et al., 2025), (Liu et al., 2023), (Hu et al., 2025), (Fouda et al., 2026), (Qiu et al., 2026), (Jaiswal et al., 2024), (Ozgun et al., 2025), (Louie et al., 2024), (Basar et al., 2024), (Yang et al., 2025b), (Kermani et al., 2025), (Basar et al., 2025), (Liu et al., 2025), (Chen et al., 2025), (Chen et al., 2023), (Schmidgall et al., 2025)

Table 3: Papers Classified by Theory Operationalization in Evaluation

Evaluation Type	Example Papers
Validated scale <i>Standardized psychological or clinical scales</i>	(Xiao et al., 2024), (Qiu et al., 2025), (Li et al., 2024), (Wang et al., 2024a), (Schmidmaier et al., 2025b), (Schmidmaier et al., 2025a), (Yang et al., 2025a), (Chen et al., 2025), (Wang et al., 2025a)
Established therapy/counseling coding system <i>Therapy process or outcome coding</i>	(Xiao et al., 2024), (Zhang et al., 2025), (Lee et al., 2024), (Feng et al., 2025), (Kim et al., 2025c), (Kim et al., 2025b), (Tanana et al., 2019)
Theory-specific rubric <i>Custom rubrics based on theoretical constructs</i>	(Xiao et al., 2024), (Gabriel et al., 2024), (Qiu et al., 2025), (Aghakhani et al., 2025), (Wang et al., 2026), (Sungarda et al., 2025), (Liu et al., 2023), (Chiu et al., 2024), (Li et al., 2024), (Wang et al., 2024a), (Fouda et al., 2026), (Cai et al., 2026), (Qiu et al., 2026), (Nguyen et al., 2025), (Zhang et al., 2025), (Kim et al., 2025a), (Xu et al., 2025a), (Ozgun et al., 2025), (Louie et al., 2026), (Louie et al., 2024), (Wang et al., 2024b), (Lee et al., 2024), (Basar et al., 2024), (Yang et al., 2025b), (Yang et al., 2025a), (Basar et al., 2025), (Feng et al., 2025), (Kim et al., 2025c), (Kim et al., 2025b), (Samuel et al., 2025), (Chen et al., 2025), (Na, 2024), (Sharma et al., 2023), (Kang et al., 2024), (Wang et al., 2025a)
Custom expert rating <i>Expert judgments on custom dimensions</i>	(Xiao et al., 2024), (Gabriel et al., 2024), (Haydarov et al., 2025), (Wang et al., 2026), (Qiu et al., 2026), (Kim et al., 2025a), (Louie et al., 2026), (Louie et al., 2024), (Wang et al., 2024b), (Basar et al., 2024), (Yang et al., 2025b), (Basar et al., 2025), (Kim et al., 2025c), (Liu et al., 2025), (Na, 2024), (Sharma et al., 2023), (Kang et al., 2024), (Schmidgall et al., 2025), (Wang et al., 2025a)
Affective/emotional dynamics metric <i>Emotion recognition or affective computing metrics</i>	(Haydarov et al., 2025), (Schmidmaier et al., 2025a), (Wang et al., 2025c)
Clinical diagnostic benchmark <i>Clinical diagnosis or screening benchmarks</i>	(Qiu et al., 2025), (Hu et al., 2025), (Fouda et al., 2026), (Cai et al., 2026), (Ozgun et al., 2025), (Liu et al., 2025), (Chen et al., 2023), (Schmidgall et al., 2025)
Theory-specific benchmark/task <i>Task-specific performance benchmarks</i>	(Qiu et al., 2025), (Aghakhani et al., 2025), (Wang et al., 2026), (Sungarda et al., 2025), (Chiu et al., 2024), (Hu et al., 2025), (Fouda et al., 2026), (Nguyen et al., 2025), (Zhang et al., 2025), (Kim et al., 2025a), (Jaiswal et al., 2024), (Xu et al., 2025a), (Ozgun et al., 2025), (Wang et al., 2024b), (Lee et al., 2024), (Kermani et al., 2025), (Feng et al., 2025), (Kim et al., 2025b), (Liu et al., 2025), (Wang et al., 2025c), (Samuel et al., 2025), (Na, 2024), (Sharma et al., 2023), (Kang et al., 2024), (Chen et al., 2023), (Tanana et al., 2019)

Table 4: Evaluation Types Used in Included Studies

Evaluator Type	Count (%)
Human Experts	34 (65.4%)
LLM Judges	30 (57.7%)
Automatic Metrics	35 (67.3%)
All Three (Expert + LLM + Auto)	15 (28.8%)

Table 5: Distribution of Evaluator Types Across Studies ($N = 52$)

Evaluator Type	Papers
Human Experts	(Aghakhani et al., 2025), (Berrezueta-Guzman et al., 2024), (Cai et al., 2026), (Chen et al., 2023), (Chen et al., 2025), (Deshpande et al., 2025), (Feng et al., 2025), (Gabriel et al., 2024), (Haydarov et al., 2025), (Kim et al., 2025b), (Kim et al., 2025a), (Kim et al., 2025c), (Kuhail et al., 2025), (Lee et al., 2024), (Li et al., 2024), (Liu et al., 2025), (Louie et al., 2024), (Louie et al., 2026), (Na, 2024), (Nguyen et al., 2025), (Qiu et al., 2026), (Samuel et al., 2025), (Schmidgall et al., 2025), (Sharma et al., 2023), (Wang et al., 2024b), (Wang et al., 2026), (Wang et al., 2025a), (Xiao et al., 2024), (Xie et al., 2025), (Xu et al., 2025a), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b), (Zhang et al., 2025)
LLM Judges	(Aghakhani et al., 2025), (Cai et al., 2026), (Chiu et al., 2024), (Feng et al., 2025), (Fouda et al., 2026), (Haydarov et al., 2025), (Kim et al., 2025b), (Kim et al., 2025c), (Lee et al., 2024), (Li et al., 2024), (Liu et al., 2023), (Liu et al., 2025), (Louie et al., 2026), (Nguyen et al., 2025), (Ozgun et al., 2025), (Qiu et al., 2025), (Qiu et al., 2026), (Samuel et al., 2025), (Schmidgall et al., 2025), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2025b), (Wang et al., 2026), (Wang et al., 2025a), (Xiao et al., 2024), (Xie et al., 2025), (Xu et al., 2025a), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b)
Automatic Metrics	(Aghakhani et al., 2025), (Basar et al., 2025), (Chen et al., 2023), (Chiu et al., 2024), (Deshpande et al., 2025), (Feng et al., 2025), (Fouda et al., 2026), (Gabriel et al., 2024), (Hu et al., 2025), (Kang et al., 2024), (Kermani et al., 2025), (Kim et al., 2025b), (Li et al., 2024), (Louie et al., 2026), (Na, 2024), (Nguyen et al., 2025), (Ozgun et al., 2025), (Qiu et al., 2025), (Samuel et al., 2025), (Schmidgall et al., 2025), (Schmidmaier et al., 2025a), (Sharma et al., 2023), (Sungarda et al., 2025), (Tanana et al., 2019), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2025b), (Wang et al., 2026), (Wang et al., 2025c), (Wang et al., 2025a), (Xie et al., 2025), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b), (Zhang et al., 2025)
All Three (Expert + LLM + Auto)	(Aghakhani et al., 2025), (Feng et al., 2025), (Kim et al., 2025b), (Li et al., 2024), (Louie et al., 2026), (Nguyen et al., 2025), (Samuel et al., 2025), (Schmidgall et al., 2025), (Wang et al., 2024b), (Wang et al., 2026), (Wang et al., 2025a), (Xie et al., 2025), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b)

Table 6: Papers by Evaluator Type

Has Rubric	Reliability Reported	Count
Yes	Yes	13
Yes	No	31
No	Yes	2
No	No	3

Table 7: Scoring Rubric and Inter-Rater Reliability Reporting

Category	Papers
Rubric=Yes, Reliability=Yes	(Aghakhani et al., 2025), (Cai et al., 2026), (Chiu et al., 2024), (Feng et al., 2025), (Gabriel et al., 2024), (Kuhail et al., 2025), (Li et al., 2024), (Louie et al., 2024), (Samuel et al., 2025), (Xie et al., 2025), (Xu et al., 2025a), (Yang et al., 2025a), (Yang et al., 2025b)
Rubric=Yes, Reliability=No	(Basar et al., 2024), (Basar et al., 2025), (Berrezueta-Guzman et al., 2024), (Chen et al., 2023), (Chen et al., 2025), (Deshpande et al., 2025), (Fouda et al., 2026), (Kang et al., 2024), (Kim et al., 2025b), (Kim et al., 2025a), (Kim et al., 2025c), (Lee et al., 2024), (Liu et al., 2023), (Liu et al., 2025), (Louie et al., 2026), (Na, 2024), (Nguyen et al., 2025), (Ozgun et al., 2025), (Qiu et al., 2025), (Qiu et al., 2026), (Schmidgall et al., 2025), (Schmidmaier et al., 2025b), (Sharma et al., 2023), (Sungarda et al., 2025), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2026), (Wang et al., 2025a), (Xiao et al., 2024), (Xu et al., 2025b), (Zhang et al., 2025)
Rubric=No, Reliability=Yes	(Haydarov et al., 2025), (Wang et al., 2025b)
Rubric=No, Reliability=No	(Jaiswal et al., 2024), (Kermani et al., 2025), (Schmidmaier et al., 2025a)

Table 8: Papers by Rubric Use and Reliability Reporting

Evaluation Dimension	Count (%)
Utility	44 (84.6%)
Fidelity	33 (63.5%)
Realism	29 (55.8%)
Emotional Plausibility	28 (53.8%)
Consistency	20 (38.5%)
Safety	14 (26.9%)
Human Learning / Outcomes	10 (19.2%)

Table 9: Evaluation Dimension Coverage Across Studies ($N = 52$)

Evaluation Dimension	Papers
Utility	(Aghakhani et al., 2025), (Basar et al., 2024), (Basar et al., 2025), (Berrezueta-Guzman et al., 2024), (Cai et al., 2026), (Chen et al., 2023), (Chen et al., 2025), (Deshpande et al., 2025), (Feng et al., 2025), (Fouda et al., 2026), (Gabriel et al., 2024), (Haydarov et al., 2025), (Hu et al., 2025), (Jaiswal et al., 2024), (Kang et al., 2024), (Kermani et al., 2025), (Kim et al., 2025b), (Kim et al., 2025a), (Kim et al., 2025c), (Kuhail et al., 2025), (Lee et al., 2024), (Li et al., 2024), (Liu et al., 2023), (Louie et al., 2024), (Louie et al., 2026), (Na, 2024), (Nguyen et al., 2025), (Ozgun et al., 2025), (Qiu et al., 2025), (Qiu et al., 2026), (Schmidgall et al., 2025), (Schmidmaier et al., 2025b), (Sharma et al., 2023), (Tanana et al., 2019), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2026), (Wang et al., 2025a), (Xiao et al., 2024), (Xie et al., 2025), (Xu et al., 2025a), (Xu et al., 2025b), (Yang et al., 2025a), (Zhang et al., 2025)
Fidelity	(Chen et al., 2023), (Chen et al., 2025), (Chiu et al., 2024), (Deshpande et al., 2025), (Feng et al., 2025), (Fouda et al., 2026), (Gabriel et al., 2024), (Kang et al., 2024), (Kim et al., 2025b), (Kim et al., 2025a), (Kim et al., 2025c), (Li et al., 2024), (Liu et al., 2023), (Liu et al., 2025), (Louie et al., 2024), (Louie et al., 2026), (Na, 2024), (Ozgun et al., 2025), (Qiu et al., 2026), (Samuel et al., 2025), (Schmidgall et al., 2025), (Tanana et al., 2019), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2025b), (Wang et al., 2026), (Wang et al., 2025c), (Wang et al., 2025a), (Xiao et al., 2024), (Xie et al., 2025), (Yang et al., 2025a), (Yang et al., 2025b), (Zhang et al., 2025)
Realism	(Basar et al., 2024), (Basar et al., 2025), (Berrezueta-Guzman et al., 2024), (Chen et al., 2023), (Chen et al., 2025), (Chiu et al., 2024), (Deshpande et al., 2025), (Kim et al., 2025b), (Kim et al., 2025a), (Kuhail et al., 2025), (Lee et al., 2024), (Liu et al., 2023), (Liu et al., 2025), (Louie et al., 2024), (Louie et al., 2026), (Ozgun et al., 2025), (Qiu et al., 2026), (Samuel et al., 2025), (Schmidgall et al., 2025), (Tanana et al., 2019), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2025b), (Wang et al., 2026), (Wang et al., 2025c), (Xie et al., 2025), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b)
Emotional Plausibility	(Berrezueta-Guzman et al., 2024), (Cai et al., 2026), (Chen et al., 2023), (Chen et al., 2025), (Gabriel et al., 2024), (Haydarov et al., 2025), (Kang et al., 2024), (Kim et al., 2025b), (Kim et al., 2025c), (Kuhail et al., 2025), (Liu et al., 2023), (Liu et al., 2025), (Ozgun et al., 2025), (Qiu et al., 2026), (Schmidgall et al., 2025), (Schmidmaier et al., 2025b), (Sharma et al., 2023), (Wang et al., 2024b), (Wang et al., 2024a), (Wang et al., 2026), (Wang et al., 2025c), (Wang et al., 2025a), (Xiao et al., 2024), (Xie et al., 2025), (Xu et al., 2025a), (Xu et al., 2025b), (Yang et al., 2025a), (Zhang et al., 2025)
Consistency	(Berrezueta-Guzman et al., 2024), (Chen et al., 2023), (Chen et al., 2025), (Chiu et al., 2024), (Fouda et al., 2026), (Gabriel et al., 2024), (Haydarov et al., 2025), (Kim et al., 2025a), (Li et al., 2024), (Liu et al., 2025), (Louie et al., 2024), (Ozgun et al., 2025), (Qiu et al., 2026), (Wang et al., 2024a), (Wang et al., 2025b), (Wang et al., 2026), (Wang et al., 2025a), (Xiao et al., 2024), (Yang et al., 2025a), (Yang et al., 2025b)
Safety	(Basar et al., 2024), (Basar et al., 2025), (Berrezueta-Guzman et al., 2024), (Cai et al., 2026), (Fouda et al., 2026), (Gabriel et al., 2024), (Li et al., 2024), (Qiu et al., 2025), (Samuel et al., 2025), (Sharma et al., 2023), (Wang et al., 2024a), (Xiao et al., 2024), (Xie et al., 2025), (Xu et al., 2025a)
Human Learning / Outcomes	(Berrezueta-Guzman et al., 2024), (Chen et al., 2025), (Feng et al., 2025), (Li et al., 2024), (Louie et al., 2026), (Qiu et al., 2026), (Tanana et al., 2019), (Wang et al., 2024b), (Wang et al., 2025a), (Xiao et al., 2024)

Table 10: Papers by Evaluation Dimension

Interaction Level	Static	Pattern-level	Dynamic	(unspecified)
Single-turn	8	0	0	4
Short Multi-turn	1	0	0	4
Extended Dialogue	16	7	7	1
Longitudinal	0	0	2	0

Table 11: Interaction Level by Behavioral Evaluation Depth ($N = 52$)

Category	Papers
Single-turn / Static	(Berrezueta-Guzman et al., 2024), (Fouda et al., 2026), (Gabriel et al., 2024), (Hu et al., 2025), (Liu et al., 2023), (Samuel et al., 2025), (Xu et al., 2025a), (Zhang et al., 2025)
Single-turn / (unspecified)	(Haydarov et al., 2025), (Na, 2024), (Nguyen et al., 2025), (Sharma et al., 2023)
Short Multi-turn / Static	(Basar et al., 2024)
Short Multi-turn / (unspecified)	(Basar et al., 2025), (Cai et al., 2026), (Jaiswal et al., 2024), (Kuhail et al., 2025)
Extended Dialogue / Static	(Chen et al., 2023), (Kang et al., 2024), (Kim et al., 2025a), (Kim et al., 2025c), (Lee et al., 2024), (Liu et al., 2025), (Louie et al., 2024), (Ozgun et al., 2025), (Qiu et al., 2025), (Qiu et al., 2026), (Schmidgall et al., 2025), (Wang et al., 2024b), (Wang et al., 2025b), (Wang et al., 2026), (Xiao et al., 2024), (Xie et al., 2025)
Extended Dialogue / Pattern-level	(Chen et al., 2025), (Kim et al., 2025b), (Schmidmaier et al., 2025b), (Schmidmaier et al., 2025a), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b)
Extended Dialogue / Dynamic	(Chiu et al., 2024), (Feng et al., 2025), (Louie et al., 2026), (Sungarda et al., 2025), (Tanana et al., 2019), (Wang et al., 2024a), (Wang et al., 2025c)
Extended Dialogue / (unspecified)	(Deshpande et al., 2025)
Longitudinal / Dynamic	(Li et al., 2024), (Wang et al., 2025a)

Table 12: Papers by Interaction Level and Behavioral Evaluation Depth

Human Learning	Lay Users	User Study	Count
Yes	Yes	Yes	4
Yes	No	Yes	5
Yes	No	No	1
No	Yes	Yes	5
No	Yes	No	4
No	No	Yes	3
No	No	No	30

Table 13: Human Learning Outcomes by Lay User Involvement and User Study Design ($N = 52$)

Category	Papers
HL=Yes, LU=Yes, US=Yes	(Chen et al., 2025), (Feng et al., 2025), (Louie et al., 2026), (Tanana et al., 2019)
HL=Yes, LU=No, US=Yes	(Berrezueta-Guzman et al., 2024), (Li et al., 2024), (Qiu et al., 2026), (Wang et al., 2024b), (Xiao et al., 2024)
HL=Yes, LU=No, US=No	(Wang et al., 2025a)
HL=No, LU=Yes, US=Yes	(Chen et al., 2023), (Jaiswal et al., 2024), (Schmidmaier et al., 2025b), (Schmidmaier et al., 2025a), (Sharma et al., 2023)
HL=No, LU=Yes, US=No	(Basar et al., 2024), (Basar et al., 2025), (Kang et al., 2024), (Lee et al., 2024)
HL=No, LU=No, US=Yes	(Kim et al., 2025a), (Liu et al., 2025), (Louie et al., 2024)
HL=No, LU=No, US=No	(Aghakhani et al., 2025), (Cai et al., 2026), (Chiu et al., 2024), (Deshpande et al., 2025), (Fouda et al., 2026), (Gabriel et al., 2024), (Haydarov et al., 2025), (Hu et al., 2025), (Kermani et al., 2025), (Kim et al., 2025b), (Kim et al., 2025c), (Kuhail et al., 2025), (Liu et al., 2023), (Na, 2024), (Nguyen et al., 2025), (Ozgun et al., 2025), (Qiu et al., 2025), (Samuel et al., 2025), (Schmidgall et al., 2025), (Sungarda et al., 2025), (Wang et al., 2024a), (Wang et al., 2025b), (Wang et al., 2026), (Wang et al., 2025c), (Xie et al., 2025), (Xu et al., 2025a), (Xu et al., 2025b), (Yang et al., 2025a), (Yang et al., 2025b), (Zhang et al., 2025)

Table 14: Papers by Human Learning, Lay Users, and User Study